

基于 PaddleOCR-VL 的电路原理图 OCR 研究与实现

数据集构建、模型训练与多维度评估

张健宁

中山大学 — Sun Yat-sen University

<https://github.com/ZhangJ83/circuit-ocr-paddle>

2026 年 7 月 20 日

Dataset: github.com/ZhangJ83/circuit_ocr_dataset_final

Outline

- 1 Dataset Construction & Evaluation
- 2 Model & Training
- 3 Experiments & Results
- 4 Key Findings & Future Work
- 5 Summary

Dataset: Four-Dimension Evaluation Framework

Dimension	Score	Weight	Key Evidence
1.1 Data Scale	4/5	5	1520 train + 150 test, 4810 OCR instances
1.2 Annotation Acc.	4.5/5	5	7-round verification, quality report w/ visuals
1.3 Data Diversity	2.5/5	5	3 sources, 20 real photos (1.3%)
1.4 Difficulty	4.5/5	5	OCR-based + 3 visual dimensions
Total	15.5	/20	

Dataset repo: `circuit_ocr_dataset_final` — Quality report: `quality_report/`

1.1 Data Scale (4/5)

Statistics

- **Training:** 1,520 samples (1200 schematics + 300 synth text + 20 real photos)
- **Test:** 150 samples (held-out)
- **Validation:** 150 samples
- **Total OCR instances:** 4,810 (test set alone)
- **10 component types:** R, C, D, U, J, Q, L, LED, F, Y

Evaluation

- ✓ Test samples ≥ 100 (150)
- ✓ OCR instances ≥ 1000 (4810)
- ✓ No copyright issues (self-drawn KiCad)
- ✗ Test set 150 is adequate but could reach 500+

4/5 — 规模达标

1.2 Annotation Accuracy (4.5/5)

7-Round Verification Pipeline

- 1 JSON Schema validation
- 2 Image path verification
- 3 Control character scan
- 4 Label format regex
- 5 Unit normalization
- 6 KiCad netlist cross-check
- 7 Manual spot-check (10%)

Quality Metrics (from Report)

- Component labels: **99%**
- Parameter values: **97%**
- Pin numbers: **98%**
- Text completeness: **95%**
- Illegal characters: **0.1%**

4.5/5 — 质量报告 + 12 组可视化对比

1.2 Annotation Quality: Visual Spot-Check

`../quality_report/visual_check/sample_01.png`

1.3 Data Diversity (2.5/5) — Weakest Dimension

Current State

- **3 data sources:** KiCad exports (1200), synthetic text (300), real photos (20)
- Real photos introduce **natural noise:** lighting, perspective, CMOS artifacts
- 10 component types covered
- Label length: 203–5,749 chars

Gaps

- ✗ All schematics from **single engine** (KiCad)
- ✗ Real photos only **1.3%** of training data
- ✗ No scanned documents, blur, occlusion
- ✗ No PCB photos, no datasheet screenshots

2.5/5 — 优先改进方向

1.4 Difficulty Rationality (4.5/5)

OCR-Based Difficulty (Original)

Level	OCR Instances	Ratio
Easy	15	20%
Medium	15–35	40%
Hard	35	40%

✓ 2:4:4

ratio matches real distribution

Multi-Dimensional Labels (NEW)

- **Text Density:** Low / Medium / High
- **Structure:** Simple / Moderate / Complex
- **Value Richness:** Sparse / Moderate / Dense
- **Composite Visual Score:** 0–3

✓ All 150 test samples have difficulty field

4.5/5 — OCR + 3 visual dimensions

LoRA Config


- $r = 16$, $\alpha = 32$, dropout
- Q/K/V/O + MLP layers
- 5.73M / 0.9B (0.6%)
- Projector frozen

Training

- AdamW, weight_decay=0.1
- Warmup + Cosine LR
- GradAccum=4, bfloat16
- RTX 4060 8GB

Training Pipeline & Data Flow

Training Pipeline Overview



Data Collection → Training → Evaluation

KiCad SVG → stroked-text extraction → 7-round verification →
+ synth text (20%) + real photos → PaddleOCR-VL LoRA → CompF1 + JointF1 + NED

Phase 1: Four-Experiment Design

Exp	Resolution	LR	Dropout	Epochs	Special	CompF1
exp1	384px	2×10^{-5}	0.05	2	Baseline	0.037
exp2	512px	2×10^{-5}	0.05	2	HiRes	0.058
exp3	384px	1×10^{-5}	0.10	3	Anti-Overfit	0.028
exp4	384px	2×10^{-5}	0.05	2	Unfrozen Proj.	0.092

Data: 1200 clean schematics — All models below V10-Fixed baseline (CompF1=0.206, 3800 samples)

Modality Collapse Problem

Collapse Pattern

- **S200**: “1,2,3,4,5...” (pure counting)
- **S400**: “+3.3V, +3.3V...” (voltage repetition)
- **S600**: Partial recovery
- Small dataset → model ignores image

Root Cause

- 1200 samples insufficient for VLM
- Model learns shortcut: output frequent patterns
- Training loss ↓ but quality ↓
- Loss-quality mismatch

Solution: Inject synthetic text-image pairs to force visual-text alignment


Phase 2: Synthetic Data Anti-Collapse Strategy

Method


- 1 Generate 300 synthetic text images (6 templates)
- 2 Mix at 20% ratio with circuit data
- 3 Total: 1500 training samples
- 4 Two configurations tested:
 - exp5: Anti-Overfit ($lr=1e-5$, $do=0.1$)
 - exp6: Baseline ($lr=2e-5$, $do=0.05$)

Results

Synthetic Data Ablation: 300 Text-Image Pairs Mixed at 20%



All Experiments: CompF1 Comparison



Phase 2 (exp5/exp6) with synthetic data: $4.2\times$ improvement over Phase 1

Training Quality Evolution



Key Insight

Without synthetic data (exp3): gradual, peaks at S800. With synthetic data (exp5): **immediate high quality at S200**, stable through S800.

Comprehensive Results

Metric	Phase 1 Best	Phase 2 Best	V10-Fixed S600	Target
CompF1	0.092 (exp4)	0.156 (exp6)	0.206 (3800 samples)	≥0.25
JointF1	0.005 (exp1/4)	0.011 (exp5/6)	0.019	≥0.05
NED	0.941 (exp4)	0.941 (exp5)	0.803	≥0.75

Status

- Synthetic data strategy **confirmed effective**: CompF1 ↑4.2×, JointF1 ↑2.2×
- Model remains below usable threshold (JointF1 0.011 vs target 0.05)
- **Bottleneck**: 1200 training schematics insufficient; need 3000+ for generalization

Key Findings

What Worked (✓)

- 1 **Synthetic data mixing** — $4.2\times$ CompF1 gain
- 2 **Anti-overfit config** — stable training, no degradation
- 3 **Frozen projector + LoRA** — prevents catastrophic collapse
- 4 **Multi-metric evaluation** — CompF1 + JointF1 + NED + RepRate
- 5 **7-round annotation pipeline** — $\sim 99\%$ label accuracy

What Didn't Work (✗)

- 1 **512px high-res** — no benefit over 384px
- 2 **Unfrozen projector** — higher CompF1 but still collapsed
- 3 **Loss-based checkpoint selection** — loss $\neq$ quality
- 4 **1200 samples** — insufficient for JointF1 > 0.05
- 5 **Single CAD source** — limits diversity score (1.3: 2.5/5)

Core insight: Synthetic data teaches visual-text alignment. Real schematic diversity needed for generalization.

Dataset Evaluation Summary

Dimension	Score	Key Strength	Key Gap
1.1 Scale	4.0/5	4810 OCR instances	Test set could be larger
1.2 Accuracy	4.5/5	Quality report + visuals	No dual-annotator Kappa
1.3 Diversity	2.5/5	3 sources, real photos	Only 20 photos, single CAD
1.4 Difficulty	4.5/5	OCR + 3 visual dims	No visual degradation labels
Total	15.5/20		Priority: ↑1.3 Diversity

Full report: `circuit_ocr_dataset_final`

Limitations & Future Work

Current Limitations

- JointF1 (0.011) $\ll$ usable threshold (0.05)
- 1200 schematics from single source (KiCad)
- Only 20 real-camera photos
- No degraded samples (blur, tilt, low-light)
- Validation set (3 samples) too small

Next Steps

- 1 **Scale data:** 3000+ schematics from multiple EDA tools
- 2 **Real photos:** 200+ photographed circuits
- 3 **Degraded samples:** blur, tilt, low-light variants
- 4 **Two-stage training:** synth pre-train $\rightarrow$ real fine-tune
- 5 **Larger LoRA:** try $r = 32$, $r = 64$

Summary

Project Achievements

- Built **high-quality dataset** (15.5/20 score, 7-round verification)
- Established **complete training pipeline** with anti-collapse strategy
- Achieved **CompF1 = 0.156** via synthetic data mixing (Phase 2)
- Generated **quality report** with visual spot-checks + multi-dim difficulty labels

Next Milestone

- ① Expand diversity (1.3): 200+ real photos, multi-CAD, degraded samples
- ② Scale training data: 3000+ schematics
- ③ Target: **JointF1 $\geq$ 0.05, CompF1 $\geq$ 0.25**

`github.com/ZhangJ83/circuit-ocr-paddle` — `circuit_ocr_dataset_final`

Thank You!

谢谢!

张健宁 — 中山大学

`github.com/ZhangJ83/circuit-ocr-paddle`

`github.com/ZhangJ83/circuit_ocr_dataset_final`

PaddleOCR-VL-0.9B + LoRA — Dataset Score: 15.5/20